

---

# Q-Learning: Deep Deterministic Policy Gradient

Zander Zemliak  
CMSI 5350  
December 6, 2022

---

# Overview of Topics

- I. Introduction to Reinforcement Learning (RL)
  - II. RL algorithm: Q-Learning
  - III. Deep Q-Learning
  - IV. DQN Algorithm: Deep Deterministic Policy Gradient (DDPG)
  - V. Stock Trading Agent built Using the DDPG Algorithm
-

---

# Reinforcement Learning

---

---


# Introduction to RL

- RL is one of the main disciplines of machine learning.
  - Differences between the three main ML disciplines:
    - Supervised learning is concerned with learning an approximate mapping between an input and an output using a labeled dataset.
    - Unsupervised learning is concerned with finding a pattern in unlabeled data.
    - RL agents learn by interacting with their environment.
  - RL agents main objective is to take actions that maximize reward.
  - A dog learns that when it sits on command it gets a treat.
-

---

# Markov Decision Process (MDP)

- MDP is an essential framework for RL, it is a mathematical model of decision making where the outcomes can be partly arbitrary or under control.
- Four main components for MDP: states, actions, effects of actions on future states, effects of actions on future rewards.



- Agriculture ex: How much to plant given the current state of water and soil?

---

## Brief MDP Math

- An MDP can be described as:  $MDP = (S, A, T, R, \gamma)$ 
  - $S$  = set of possible states
  - $A$  = set of possible actions
  - $T$  = set of transition probabilities where the probability of going from  $s$  to  $s'$  is described as  $P(s'|s, a)$ .
  - $R$  = set of rewards
  - $\gamma$  = discount factor (determines how much agent cares about future rewards and is necessary for algorithm convergence)

- We will use the Bellman equation to evaluate the max current and future reward:

$$V(s) = \max_a (R(s, a) + \gamma V(s'))$$

- The value for our current state,  $V(s)$ , is equal to the action which maximizes current reward in state  $s$  with action  $a$  plus the value of the next state  $s'$ .
-

---

# Q-Learning

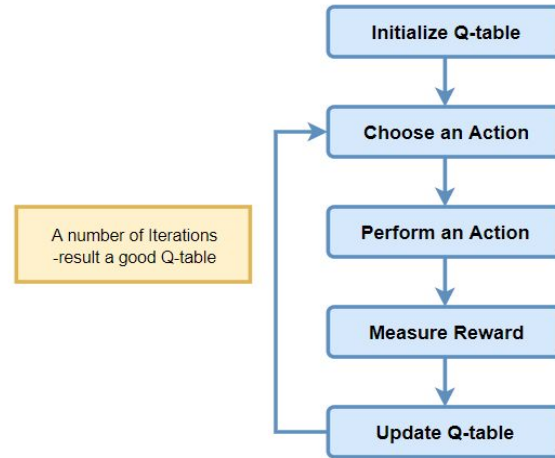
---

# Q-Learning

- Q-learning is a *value based* RL algorithm that attempts to learn the quality (Q) of actions in particular states that are defined by future value.
- The mapping of action to state is known as the policy ( $\pi$ ) and these values are stored in a Q-table with two columns: state-action pair and value.
- Agents choose actions in certain states that result in the highest reward  $Q^*(s,a)$ .
- Q-table values are updated using TD:

$$Q_{\text{new}}(s, a) = Q(s, a) + \alpha [R(s,a) + \gamma \max_a Q(s', a=a) - Q(s, a)]$$

- TD target looks very familiar to Bellman equation...
- This will converge towards  $\pi^*$





---


# Deep Q-Learning

---

---

# Deep Q-Learning

- Very similar idea to Q-Learning but instead of using a Q-table with (state, action) pairs mapped to a Q-value for choosing actions, we have a neural network that maps input states to (action, Q-value) pairs.
- Since in real life we could have many (state, action) pairs for a given environment, it is better to use a NN for a large state space because of its generalizability.



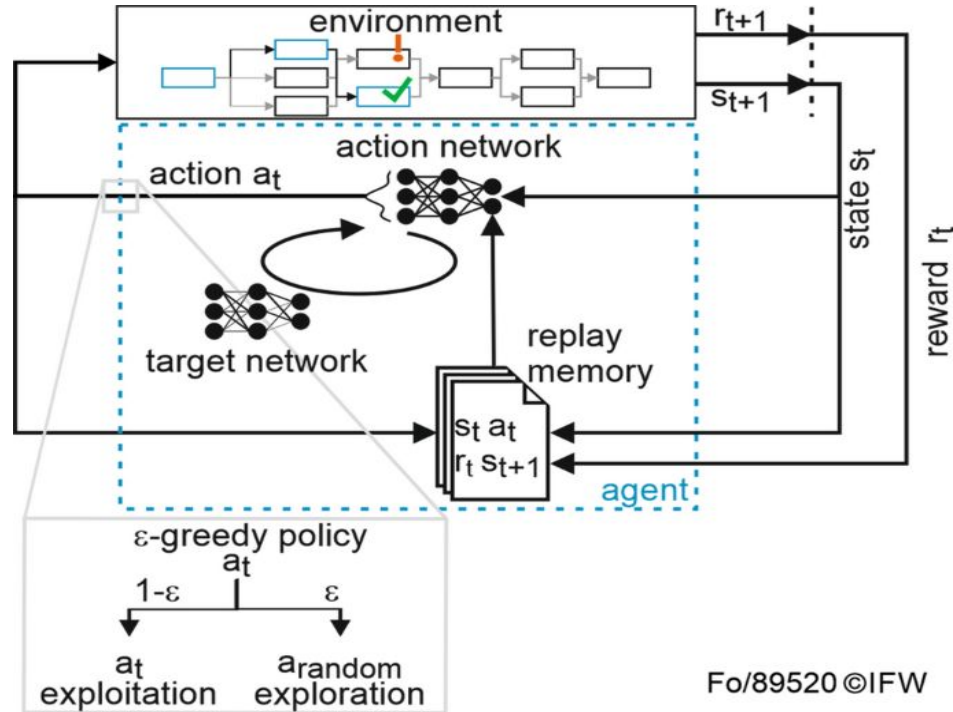
---

# Training a DQN

- DQN training works by taking an action via epsilon-greedy, observing the episode of that action, storing it in memory, and then sampling from the replay memory to update main prediction NN.
- We use two networks: a main and target network
  - Main NN: Where our Q-values are derived
  - Target NN: Same parameters as the main NN however, this network is not trained. It is synchronized with the main NN after N steps. This allows us to train the main NN with a single target thus making training stable.
- Weights are updated in the main network via the Q-Learning Loss function:

$$L(\theta) = \left( \overbrace{(r + \gamma \max_{a_{t+1}} Q(s_{t+1}, a_{t+1}; \theta^{target}))}^{\text{Target Q value}} - \underbrace{Q(s, a; \theta^{pred})}_{\text{Predicted Q value}} \right)^2$$

# Training a DQN



Fo/89520 ©IFW

---

# **DQN Algorithm: Deep Deterministic Policy Gradient (DDPG)**

---

---

# DDPG Overview

- Q-Learning and Deep Q-Learning are great for environments that have discrete action-state spaces. But, what if our action-state space is *continuous*?
- What if the amount of actions we can take in any given state is not finite?
- Remember, the ultimate goal of our agent is to choose the best action for any state:

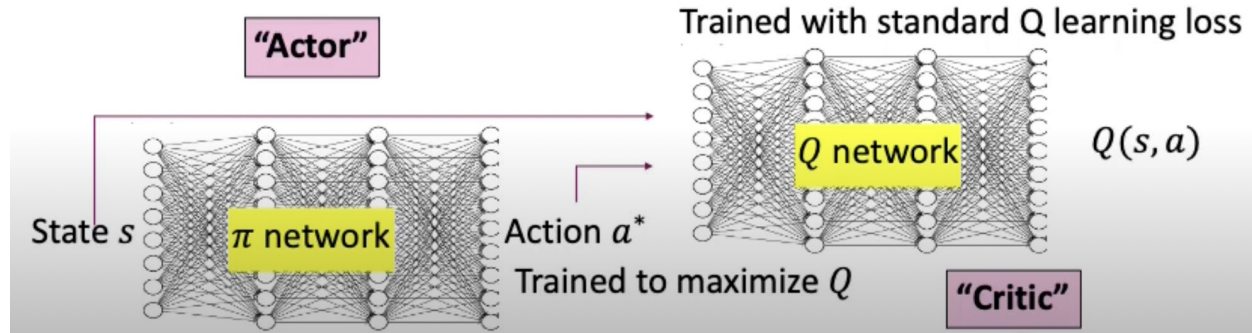
$$\pi(s) = a^*(s)$$

$$\text{where; } a^*(s) = \operatorname{argmax}_a Q(s, a)$$

- To solve this issue of having a continuous action-state space, DDPG uses an actor network and critic network.
    - Actor network: Responsible for mapping state to  $a^*$
    - Critic network: Responsible for mapping  $(s, a^*)$  tuple to  $Q(s, a)$
-

# DDPG Training

- The actor network is trained to maximize the Q-value from the critic network so its update rule uses gradient **ascent**.
- The critic network is updated using the traditional Q-Learning Loss function with.
- Similar to DQN, each actor and critic network both have target networks to help with stabilizing learning.



---

# DDPG Algorithm Steps

1. Initialize actor  $\theta$  params ( $\pi$  net), critic  $\phi$  params (Q net), and empty replay buffer B
2. Set target params equal to their respective main networks (just like DQ Learning)
3. Repeat until convergence:
  - a. Observe state  $s$  and get  $a^*$  from critic and execute  $a^*$
  - b. Observe state  $s'$ , get reward  $r$ , and set  $d$  = is in terminal state (0=no, 1=yes)
  - c. Store episode  $(s, a, r, s', d)$  in replay buffer
  - d. If  $d = 1$  reset environment parameters otherwise continue
  - e. If time to update:
    - i. For however many updates:
      1. Get sample batch from B
      2. Compute target via Bellman:  $y(r, s', d) = r + \gamma(1-d)Q_{\phi_{\text{target}}}(s', \pi_{\theta_{\text{target}}})$
      3. Update critic net via one step of GD using  $y(r, s', d)$
      4. Update actor net via one step of GA using:  $Q_{\phi}(s, \pi_{\theta}(s))$
      5. Update target networks with their respective main network



---

# **DDPG Example: Stock Trading Agent**

---

---

# Problem Justification

- Managing a portfolio of stocks to create a comfortable return can be very difficult:
    - Emotions can impact one's due diligence on a planned strategy.
    - Trading in live time can be stressful.
    - You don't know what you're doing.
  - Because the stock market changes every second of everyday, the state space is **extremely** large. DDPG is intended for large state spaces.
  - Using DDPG, we can potentially find a viable trading strategy within the vast complexities of the financial markets.
-

---

# Problem Overview

## Some Stuff for Q-Learning:


- Actions: buy, sell, or hold a stock.
- State space: N stocks we are interested in that have features such as closing, high, low, volume, day of week, rsi, ema, sma, etc..
- Reward: Return of stock of trade.
- Therefore, overall goal of agent is to maximize trading returns

## Data:

- We will have our agent only trade Apple stock (AAPL).
  - We will compare the performance of our agent to how Apple stock performed on its own in the same time frame.
  - For training, we will use 10 years of historical data from Apple.
  - For testing, we will use 8 months of historical data of Apple.
-

---

# Results



August 2021 to April 2022:

- DDPG result: +4.47%
- AAPL result: +22.7%

---

# Results

- Model performed poorly with respect to the overall movement of Apple in the same time frame.
  - An individual buying and holding the stock with X amount of starting cash would have outperformed the trading agent significantly.
  - The model appeared to have adopted a FOMO strategy. Buying and selling with respect to Apple's price movement. Were the technical indicator features just noise?
  - A positive return appears as a success. Should you let this model trade with real money? **NO!** A positive backtest result does not indicate any kind of model generality for trading.
-

---

# Future Improvements

- Adjusting hyperparameters:
    - Changing learning rate
    - Lowering batch sample size
    - Lowering the number of steps.
  - Using different features for training (deeper research into technical indicators).
  - Using a different set of stocks that have a lot more volatility.
  - Hardcoding functions that step into the agent's actions if a certain threshold of success/failure is met. E.g. stop trading if cumulative return is +/- 5%.
-

---

**Thank you!**  
**Questions?**

---